

Домашнее задание

Sqoop:

1. Получаем список таблиц, которые созданы в postgres (нужно правильно указать ip, сейчас указан localhost)

```
sqoop list-tables --connect jdbc:postgresql://localhost/pg_db --username exporter --password exporter_pass
```

```
[student1_13@bigdataanalytics-head-0 ~]$ sqoop list-tables --connect jdbc:postgresql://185.241.193.174/pg_db --username exporter --password exporter_pass
Warning: /usr/hdp/3.1.4.0-315/accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hadoop/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/phoenix/phoenix-5.0.0.3.1.4.0-315-server.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
21/02/18 22:24:43 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7.3.1.4.0-315
21/02/18 22:24:43 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
21/02/18 22:24:43 INFO manager.SqlManager: Using default fetchSize of 1000
```

```
[student1_13@bigdataanalytics-head-0 ~]$ sqoop list-tables --connect jdbc:postgresql://student1_13@185.241.193.174/pg_db --username exporter --password exporter_pass
Warning: /usr/hdp/3.1.4.0-315/accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hadoop/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/phoenix/phoenix-5.0.0.3.1.4.0-315-server.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
21/02/18 22:41:28 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7.3.1.4.0-315
```

```
21/02/18 22:41:28 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
21/02/18 22:41:28 INFO manager.SqlManager: Using default fetchSize of
1000
21/02/18 22:41:29 ERROR manager.CatalogQueryManager: Failed to list
tables
org.postgresql.util.PSQLException: The connection attempt failed.
    at
org.postgresql.core.v3.ConnectionFactoryImpl.openConnectionImpl(Connecti
onFactoryImpl.java:262)
    at
org.postgresql.core.ConnectionFactory.openConnection(ConnectionFactory.j
ava:66)
    at org.postgresql.jdbc.PgConnection.<init>(PgConnection.java:215)
    at org.postgresql.Driver.makeConnection(Driver.java:406)
    at org.postgresql.Driver.connect(Driver.java:274)
    at java.sql.DriverManager.getConnection(DriverManager.java:664)
    at java.sql.DriverManager.getConnection(DriverManager.java:247)
    at
org.apache.sqoop.manager.SqlManager.makeConnection(SqlManager.java:903)
    at
org.apache.sqoop.manager.GenericJdbcManager.getConnection(GenericJdbcMan
ager.java:59)
    at
org.apache.sqoop.manager.CatalogQueryManager.listTables(CatalogQueryMana
ger.java:102)
    at org.apache.sqoop.tool.ListTablesTool.run(ListTablesTool.java:49)
    at org.apache.sqoop.Sqoop.run(Sqoop.java:151)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)
    at org.apache.sqoop.Sqoop.runSqoop(Sqoop.java:187)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:241)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:250)
    at org.apache.sqoop.Sqoop.main(Sqoop.java:259)
Caused by: java.net.UnknownHostException: student1_13@185.241.193.174
    at
java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:18
4)
    at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
    at java.net.Socket.connect(Socket.java:589)
    at org.postgresql.core.PGStream.<init>(PGStream.java:64)
    at
org.postgresql.core.v3.ConnectionFactoryImpl.openConnectionImpl(Connecti
onFactoryImpl.java:133)
    ... 16 more
21/02/18 22:41:29 ERROR sqoop.Sqoop: Got exception running Sqoop:
java.lang.RuntimeException: org.postgresql.util.PSQLException: The
connection attempt failed.
java.lang.RuntimeException: org.postgresql.util.PSQLException: The
connection attempt failed.
    at
org.apache.sqoop.manager.CatalogQueryManager.listTables(CatalogQueryMana
ger.java:118)
    at org.apache.sqoop.tool.ListTablesTool.run(ListTablesTool.java:49)
    at org.apache.sqoop.Sqoop.run(Sqoop.java:151)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)
    at org.apache.sqoop.Sqoop.runSqoop(Sqoop.java:187)
```

```

    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:241)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:250)
    at org.apache.sqoop.Sqoop.main(Sqoop.java:259)
Caused by: org.postgresql.util.PSQLException: The connection attempt
failed.
    at
org.postgresql.core.v3.ConnectionFactoryImpl.openConnectionImpl(Connecti
onFactoryImpl.java:262)
    at
org.postgresql.core.ConnectionFactory.openConnection(ConnectionFactory.j
ava:66)
    at org.postgresql.jdbc.PgConnection.<init>(PgConnection.java:215)
    at org.postgresql.Driver.makeConnection(Driver.java:406)
    at org.postgresql.Driver.connect(Driver.java:274)
    at java.sql.DriverManager.getConnection(DriverManager.java:664)
    at java.sql.DriverManager.getConnection(DriverManager.java:247)
    at
org.apache.sqoop.manager.SqlManager.makeConnection(SqlManager.java:903)
    at
org.apache.sqoop.manager.GenericJdbcManager.getConnection(GenericJdbcMan
ager.java:59)
    at
org.apache.sqoop.manager.CatalogQueryManager.listTables(CatalogQueryMana
ger.java:102)
    ... 7 more
Caused by: java.net.UnknownHostException: student1_13@185.241.193.174
    at
java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:18
4)
    at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
    at java.net.Socket.connect(Socket.java:589)
    at org.postgresql.core.PGStream.<init>(PGStream.java:64)
    at
org.postgresql.core.v3.ConnectionFactoryImpl.openConnectionImpl(Connecti
onFactoryImpl.java:133)
    ... 16 more

```

2. Создаем таблицу в hive, куда собираемся выгружать данные

```

CREATE TABLE lesson6_9.character(
  charid string,
  charname string,
  abbrev string, description string, speechcount int)

```

```

0: jdbc:hive2://bigdataanalytics-worker-0.nov> create database
lesson6_9;
INFO : Compiling command(queryId=hive_20210218222744_c2edda29-
af0c-4088-a449-9ae7837d34bf): create database lesson6_9
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null,
properties:null)

```

```

INFO : Completed compiling
command(queryId=hive_20210218222744_c2edda29-af0c-4088-
a449-9ae7837d34bf); Time taken: 0.017 seconds
INFO : Executing command(queryId=hive_20210218222744_c2edda29-
af0c-4088-a449-9ae7837d34bf): create database lesson6_9
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing
command(queryId=hive_20210218222744_c2edda29-af0c-4088-
a449-9ae7837d34bf); Time taken: 0.062 seconds
INFO : OK
No rows affected (0.343 seconds)
0: jdbc:hive2://bigdataanalytics-worker-0.nov> use lesson6_9;
INFO : Compiling
command(queryId=hive_20210218222759_3086689e-8a25-471f-
b5e3-88878cb68c27): use lesson6_9
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null,
properties:null)
INFO : Completed compiling
command(queryId=hive_20210218222759_3086689e-8a25-471f-
b5e3-88878cb68c27); Time taken: 0.008 seconds
INFO : Executing
command(queryId=hive_20210218222759_3086689e-8a25-471f-
b5e3-88878cb68c27): use lesson6_9
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing
command(queryId=hive_20210218222759_3086689e-8a25-471f-
b5e3-88878cb68c27); Time taken: 0.006 seconds
INFO : OK
No rows affected (0.023 seconds)
0: jdbc:hive2://bigdataanalytics-worker-0.nov> CREATE TABLE
lesson6_9.character(charid string, charname string, abbrev string,
description string, speechcount int);
INFO : Compiling
command(queryId=hive_20210218222804_e9a0ed9b-68c5-42fc-98c8-8492984142b8
): CREATE TABLE lesson6_9.character(charid string, charname string,
abbrev string, description string, speechcount int)
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null,
properties:null)
INFO : Completed compiling
command(queryId=hive_20210218222804_e9a0ed9b-68c5-42fc-98c8-8492984142b8
); Time taken: 0.027 seconds
INFO : Executing
command(queryId=hive_20210218222804_e9a0ed9b-68c5-42fc-98c8-8492984142b8
): CREATE TABLE lesson6_9.character(charid string, charname string,
abbrev string, description string, speechcount int)
INFO : Starting task [Stage-0:DDL] in serial mode
ERROR : FAILED: Execution Error, return code 1 from
org.apache.hadoop.hive.ql.exec.DDLTask. MetaException(message:Got
exception: org.apache.hadoop.security.AccessControlException Permission
denied: user=student1_13, access=EXECUTE, inode="/warehouse/tablespace/
managed/hive/lesson6_9.db":hive:hadoop:drwxrwx---
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.check(FSPermissionChecker.java:399)

```

```

    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkTraverse
(FSPermissionChecker.java:315)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:242)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:193)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkTraverse
(FSPermissionChecker.java:606)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkTraverse(FSDirec
tory.java:1806)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkTraverse(FSDirec
tory.java:1824)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.resolvePath(FSDirecto
ry.java:681)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirStatAndListingOp.getFileInfo
(FSDirStatAndListingOp.java:114)
    at
org.apache.hadoop.hdfs.server.namenode.FSNamesystem.getFileInfo(FSNamesy
stem.java:3106)
    at
org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.getFileInfo(Nam
eNodeRpcServer.java:1154)
    at
org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolServerSideTransl
atorPB.getFileInfo(ClientNamenodeProtocolServerSideTranslatorPB.java:966
)
    at
org.apache.hadoop.hdfs.protocol.proto.ClientNamenodeProtocolProtos$Clie
ntNamenodeProtocol$2.callBlockingMethod(ClientNamenodeProtocolProtos.java
)
    at
org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(P
rotobufRpcEngine.java:524)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:1025)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:876)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:822)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformatio
n.java:1730)
    at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2682)
)
INFO : Completed executing
command(queryId=hive_20210218222804_e9a0ed9b-68c5-42fc-98c8-8492984142b8
); Time taken: 0.024 seconds
Error: Error while processing statement: FAILED: Execution Error, return
code 1 from org.apache.hadoop.hive.ql.exec.DDLTask.

```

```
MetaException(message:Got exception:
org.apache.hadoop.security.AccessControlException Permission denied:
user=student1_13, access=EXECUTE, inode="/warehouse/tablesapce/managed/
hive/lesson6_9.db":hive:hadoop:drwxrwx---
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.check(FSPermi
ssionChecker.java:399)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkTraverse
(FSPermissionChecker.java:315)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:242)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:193)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkTraverse
(FSPermissionChecker.java:606)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkTraverse(FSDirec
tory.java:1806)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkTraverse(FSDirec
tory.java:1824)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.resolvePath(FSDirecto
ry.java:681)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirStatAndListingOp.getFileInfo
(FSDirStatAndListingOp.java:114)
    at
org.apache.hadoop.hdfs.server.namenode.FSNamesystem.getFileInfo(FSNamesy
stem.java:3106)
    at
org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.getFileInfo(Nam
eNodeRpcServer.java:1154)
    at
org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolServerSideTransl
atorPB.getFileInfo(ClientNamenodeProtocolServerSideTranslatorPB.java:966
)
    at
org.apache.hadoop.hdfs.protocol.proto.ClientNamenodeProtocolProtos$Clien
tNamenodeProtocol$2.callBlockingMethod(ClientNamenodeProtocolProtos.java
)
    at
org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(P
rotobufRpcEngine.java:524)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:1025)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:876)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:822)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformatio
n.java:1730)
```

```
at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2682)
) (state=08S01,code=1)
```

3. Запускаем импорт:

```
sqoop import --connect jdbc:postgresql://node3.novalocal/pg_db --
username exporter --password exporter_pass --table character --hive-
import --hive-database lesson6_9 --hive-table character
```

```
[student1_13@bigdataanalytics-head-0 ~]$ sqoop import --connect
jdbc:po//node3.novalocal/pg_db --username exporter --password
exporter_pass --table character --hive-import --hive-database lesson6_9
--hive-table character
Warning: /usr/hdp/3.1.4.0-315/accumulo does not exist! Accumulo imports
will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hadoop/lib/slf4j-
log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/phoenix/
phoenix-5.0.0.3.1.4.0-315-server.jar!/org/slf4j/impl/
StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hive/lib/log4j-
slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
21/02/18 22:35:48 INFO sqoop.Sqoop: Running Sqoop version:
1.4.7.3.1.4.0-315
21/02/18 22:35:48 WARN tool.BaseSqoopTool: Setting your password on the
command-line is insecure. Consider using -P instead.
21/02/18 22:35:48 INFO tool.BaseSqoopTool: Using Hive-specific
delimiters for output. You can override
21/02/18 22:35:48 INFO tool.BaseSqoopTool: delimiters with --fields-
terminated-by, etc.
21/02/18 22:35:48 INFO manager.SqlManager: Using default fetchSize of
1000
21/02/18 22:35:48 INFO tool.CodeGenTool: Beginning code generation
21/02/18 22:35:48 ERROR manager.SqlManager: Error executing statement:
org.postgresql.util.PSQLException: The connection attempt failed.
org.postgresql.util.PSQLException: The connection attempt failed.
at
org.postgresql.core.v3.ConnectionFactoryImpl.openConnectionImpl(Connecti
onFactoryImpl.java:262)
at
org.postgresql.core.ConnectionFactory.openConnection(ConnectionFactory.j
ava:66)
at org.postgresql.jdbc.PgConnection.<init>(PgConnection.java:215)
at org.postgresql.Driver.makeConnection(Driver.java:406)
at org.postgresql.Driver.connect(Driver.java:274)
at java.sql.DriverManager.getConnection(DriverManager.java:664)
at java.sql.DriverManager.getConnection(DriverManager.java:247)
```

```

    at
org.apache.sqoop.manager.SqlManager.makeConnection(SqlManager.java:903)
    at
org.apache.sqoop.manager.GenericJdbcManager.getConnection(GenericJdbcManager.java:59)
    at org.apache.sqoop.manager.SqlManager.execute(SqlManager.java:762)
    at org.apache.sqoop.manager.SqlManager.execute(SqlManager.java:785)
    at
org.apache.sqoop.manager.SqlManager.getColumnInfoForRawQuery(SqlManager.java:288)
    at
org.apache.sqoop.manager.SqlManager.getColumnTypesForRawQuery(SqlManager.java:259)
    at
org.apache.sqoop.manager.SqlManager.getColumnTypes(SqlManager.java:245)
    at
org.apache.sqoop.manager.ConnManager.getColumnTypes(ConnManager.java:334)
)
    at
org.apache.sqoop.orm.ClassWriter.getColumnTypes(ClassWriter.java:1879)
    at org.apache.sqoop.orm.ClassWriter.generate(ClassWriter.java:1672)
    at
org.apache.sqoop.tool.CodeGenTool.generateORM(CodeGenTool.java:106)
    at org.apache.sqoop.tool.ImportTool.importTable(ImportTool.java:524)
    at org.apache.sqoop.tool.ImportTool.run(ImportTool.java:655)
    at org.apache.sqoop.Sqoop.run(Sqoop.java:151)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)
    at org.apache.sqoop.Sqoop.runSqoop(Sqoop.java:187)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:241)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:250)
    at org.apache.sqoop.Sqoop.main(Sqoop.java:259)
Caused by: java.net.UnknownHostException: node3.novalocal
    at
java.net.AbstractPlainSocketImpl.connect(AbstractPlainSocketImpl.java:184)
    at java.net.SocksSocketImpl.connect(SocksSocketImpl.java:392)
    at java.net.Socket.connect(Socket.java:589)
    at org.postgresql.core.PGStream.<init>(PGStream.java:64)
    at
org.postgresql.core.v3.ConnectionFactoryImpl.openConnectionImpl(ConnectionFactoryImpl.java:133)
    ... 25 more
21/02/18 22:35:48 ERROR tool.ImportTool: Import failed:
java.io.IOException: No columns to generate for ClassWriter
    at org.apache.sqoop.orm.ClassWriter.generate(ClassWriter.java:1678)
    at
org.apache.sqoop.tool.CodeGenTool.generateORM(CodeGenTool.java:106)
    at org.apache.sqoop.tool.ImportTool.importTable(ImportTool.java:524)
    at org.apache.sqoop.tool.ImportTool.run(ImportTool.java:655)
    at org.apache.sqoop.Sqoop.run(Sqoop.java:151)
    at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)
    at org.apache.sqoop.Sqoop.runSqoop(Sqoop.java:187)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:241)
    at org.apache.sqoop.Sqoop.runTool(Sqoop.java:250)
    at org.apache.sqoop.Sqoop.main(Sqoop.java:259)

```


Flume:

1. Запускаем flume (нужно правильно указать пути до папки и файла) `/opt/apache-flume/bin/flume-ng agent --conf /home/admin/flm/ --conf-file /home/admin/flm/flm.conf --name StudentFlume -Dflume.root.logger=INFO,console`

```
[student1_13@bigdataanalytics-head-0 ~]$ /opt/apache-flume/bin/flume-ng
agent --conf /home/admin/flm/ --conf-file /home/admin/flm/flm.conf --
name StudentFlume -Dflume.root.logger=INFO,console
Info: Including Hadoop libraries found via (/bin/hadoop) for HDFS access
Info: Including HBASE libraries found via (/bin/hbase) for HBASE access
Info: Including Hive libraries found via () for Hive access
+ exec /usr/java/oracle-jdk/bin/java -Xmx20m
-Dflume.root.logger=INFO,console -cp '/home/admin/flm:/opt/apache-
flume/lib/*:/usr/hdp/3.1.4.0-315/hadoop/conf:/usr/hdp/3.1.4.0-315/
hadoop/lib/*:/usr/hdp/3.1.4.0-315/hadoop/*:/usr/hdp/3.1.4.0-315/
hadoop-hdfs/*:/usr/hdp/3.1.4.0-315/hadoop-hdfs/lib/*:/usr/hdp/
3.1.4.0-315/hadoop-hdfs/*:/usr/hdp/3.1.4.0-315/hadoop-mapreduce/lib/
*/usr/hdp/3.1.4.0-315/hadoop-mapreduce/*:/usr/hdp/3.1.4.0-315/
hadoop-yarn/*:/usr/hdp/3.1.4.0-315/hadoop-yarn/lib/*:/usr/hdp/
3.1.4.0-315/hadoop-yarn/*:/usr/hdp/3.1.4.0-315/tez/*:/usr/hdp/
3.1.4.0-315/tez/lib/*:/usr/hdp/3.1.4.0-315/tez/conf:/usr/hdp/
3.1.4.0-315/tez/conf_llap:/usr/hdp/3.1.4.0-315/tez/doc:/usr/hdp/
3.1.4.0-315/tez/hadoop-shim-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/
tez/hadoop-shim-2.8-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib:/
usr/hdp/3.1.4.0-315/tez/man:/usr/hdp/3.1.4.0-315/tez/tez-
api-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
common-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
dag-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
examples-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-history-
parser-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-javadoc-
tools-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-job-
analyzer-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
mapreduce-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-protobuf-
history-plugin-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
runtime-internals-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
runtime-library-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
tests-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-yarn-timeline-
cache-plugin-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-yarn-
timeline-history-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
yarn-timeline-history-with-acls-0.9.1.3.1.4.0-315.jar:/usr/hdp/
3.1.4.0-315/tez/tez-yarn-timeline-history-with-
fs-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/ui:/usr/hdp/
3.1.4.0-315/tez/lib/async-http-client-1.9.40.jar:/usr/hdp/3.1.4.0-315/
tez/lib/commons-cli-1.2.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-
codec-1.4.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-
collections-3.2.2.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-
collections4-4.1.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-io-2.4.jar:/
usr/hdp/3.1.4.0-315/tez/lib/commons-lang-2.6.jar:/usr/hdp/3.1.4.0-315/
tez/lib/commons-math3-3.1.1.jar:/usr/hdp/3.1.4.0-315/tez/lib/gcs-
```

connector-1.9.10.3.1.4.0-315-shaded.jar:/usr/hdp/3.1.4.0-315/tez/lib/
guava-28.0-jre.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-
aws-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-
azure-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-azure-
datalake-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-hdfs-
client-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-
mapreduce-client-common-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/
lib/hadoop-mapreduce-client-core-3.1.1.3.1.4.0-315.jar:/usr/hdp/
3.1.4.0-315/tez/lib/hadoop-yarn-server-timeline-
pluginstorage-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/jersey-
client-1.19.jar:/usr/hdp/3.1.4.0-315/tez/lib/jersey-json-1.19.jar:/usr/
hdp/3.1.4.0-315/tez/lib/jettison-1.3.4.jar:/usr/hdp/3.1.4.0-315/tez/lib/
jetty-server-9.3.24.v20180605.jar:/usr/hdp/3.1.4.0-315/tez/lib/jetty-
util-9.3.24.v20180605.jar:/usr/hdp/3.1.4.0-315/tez/lib/
jsr305-3.0.0.jar:/usr/hdp/3.1.4.0-315/tez/lib/metrics-core-3.1.0.jar:/
usr/hdp/3.1.4.0-315/tez/lib/protobuf-java-2.5.0.jar:/usr/hdp/
3.1.4.0-315/tez/lib/RoaringBitmap-0.4.9.jar:/usr/hdp/3.1.4.0-315/tez/
lib/servlet-api-2.5.jar:/usr/hdp/3.1.4.0-315/tez/lib/slf4j-
api-1.7.10.jar:/usr/hdp/3.1.4.0-315/tez/lib/tez.tar.gz:/usr/hdp/
3.1.4.0-315/hbase/conf:/usr/java/oracle-jdk/lib/tools.jar:/usr/hdp/
3.1.4.0-315/hbase:/usr/hdp/3.1.4.0-315/hbase/lib/animal-sniffer-
annotations-1.17.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
aopalliance-1.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/aopalliance-
repackaged-2.5.0-b32.jar:/usr/hdp/3.1.4.0-315/hbase/lib/atlas-plugin-
classloader-1.1.0.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
audience-annotations-0.5.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
avro-1.7.7.jar:/usr/hdp/3.1.4.0-315/hbase/lib/aws-java-sdk-
bundle-1.11.375.jar:/usr/hdp/3.1.4.0-315/hbase/lib/checker-
qual-2.8.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/commons-
beanutils-1.9.3.jar:/usr/hdp/3.1.4.0-315/hbase/lib/commons-cli-1.2.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/commons-codec-1.10.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/commons-collections-3.2.2.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/commons-compress-1.4.1.jar:/usr/hdp/3.1.4.0-315/
hbase/lib/commons-configuration2-2.1.1.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/commons-crypto-1.0.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/commons-
csv-1.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/commons-daemon-1.0.13.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/commons-io-2.5.jar:/usr/hdp/3.1.4.0-315/
hbase/lib/commons-lang-2.6.jar:/usr/hdp/3.1.4.0-315/hbase/lib/commons-
lang3-3.6.jar:/usr/hdp/3.1.4.0-315/hbase/lib/commons-logging-1.2.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/commons-math3-3.6.1.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/commons-net-3.6.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/curator-client-4.0.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/curator-
framework-4.0.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/curator-
recipes-4.0.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/disruptor-3.3.6.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/dnsjava-2.1.7.jar:/usr/hdp/3.1.4.0-315/
hbase/lib/ehcache-3.3.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
error_prone_annotations-2.3.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
failureaccess-1.0.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/findbugs-
annotations-1.3.9-1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/fst-2.50.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/geronimo-jcache_1.0_spec-1.0-alpha-1.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/gson-2.2.4.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/guava-28.0-jre.jar:/usr/hdp/3.1.4.0-315/hbase/lib/guice-4.0.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/guice-servlet-4.0.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/hamcrest-core-1.3.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/hbase-annotations-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/hbase-annotations-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/

hbase/lib/hbase-annotations.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-backup-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-backup.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-bridge-shim-1.1.0.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-client-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-client.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-common-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-common-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-common.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-endpoint-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-endpoint.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-examples-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-examples.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-external-blockcache-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-external-blockcache.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-hadoop2-compat-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-hadoop2-compat-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-hadoop2-compat.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-hadoop-compat-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-hadoop-compat-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-hadoop-compat.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-http-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-http.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-it-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-it-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-it.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-mapreduce-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-mapreduce-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-mapreduce.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-metrics-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-metrics-api-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-metrics-api.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-metrics.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-procedure-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-procedure.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-protocol-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-protocol.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-protocol-shaded-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-protocol-shaded.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-replication-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-replication.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-resource-bundle-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-resource-bundle.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-rest-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-rest.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-rsgroup-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-rsgroup-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-rsgroup.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-server-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-server-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-server.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-client-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-client.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-mapreduce-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-mapreduce.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-miscellaneous-2.2.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-netty-2.2.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shaded-protobuf-2.2.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-

shell-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-shell.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-spark-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-spark-it-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-spark.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-testing-util-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-testing-util.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-thrift-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-thrift.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-zookeeper-2.0.2.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-zookeeper-2.0.2.3.1.4.0-315-tests.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hbase-zookeeper.jar:/usr/hdp/3.1.4.0-315/hbase/lib/HikariCP-java7-2.4.12.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hk2-api-2.5.0-b32.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hk2-locator-2.5.0-b32.jar:/usr/hdp/3.1.4.0-315/hbase/lib/hk2-utils-2.5.0-b32.jar:/usr/hdp/3.1.4.0-315/hbase/lib/htrace-core-3.2.0-incubating.jar:/usr/hdp/3.1.4.0-315/hbase/lib/htrace-core4-4.2.0-incubating.jar:/usr/hdp/3.1.4.0-315/hbase/lib/httpclient-4.5.3.jar:/usr/hdp/3.1.4.0-315/hbase/lib/httpcore-4.4.6.jar:/usr/hdp/3.1.4.0-315/hbase/lib/j2objc-annotations-1.3.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-annotations-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-core-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-core-asl-1.9.13.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-databind-2.9.9.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-jaxrs-1.9.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-jaxrs-base-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-jaxrs-json-provider-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-mapper-asl-1.9.13.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-module-jaxb-annotations-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-module-paranamer-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-module-scala_2.11-2.9.9.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jackson-xc-1.9.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jamon-runtime-2.4.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javassist-3.20.0-GA.jar:/usr/hdp/3.1.4.0-315/hbase/lib/java-util-1.9.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.annotation-api-1.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.el-3.0.1-b08.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.inject-2.5.0-b32.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.servlet-api-3.1.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.servlet.jsp-2.3.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.servlet.jsp-api-2.3.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.servlet.jsp.jstl-1.2.0.v201105211821.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.servlet.jsp.jstl-1.2.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/javax.ws.rs-api-2.0.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jaxb-api-2.2.12.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jaxb-impl-2.2.3-1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jcip-annotations-1.0-1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jcodings-1.0.18.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jersey-client-2.25.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jersey-common-2.25.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jersey-container-servlet-core-2.25.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jersey-guava-2.25.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jersey-media-jaxb-2.25.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jersey-server-2.25.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jettison-1.3.8.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-http-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-io-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-jmx-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-jsp-9.2.28.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-schemas-3.1.M0.jar:/usr/hdp/3.1.4.0-315/

hbase/lib/jetty-security-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/
hbase/lib/jetty-server-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/jetty-servlet-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
jetty-util-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-
util-ajax-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-
webapp-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jetty-
xml-9.3.27.v20190418.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
joni-2.1.11.jar:/usr/hdp/3.1.4.0-315/hbase/lib/jsch-0.1.54.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/json-io-2.5.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
jsr311-api-1.1.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/junit-4.12.jar:/usr/
hdp/3.1.4.0-315/hbase/lib/leveldbjni-all-1.8.jar:/usr/hdp/3.1.4.0-315/
hbase/lib/libthrift-0.9.3-1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
listenablefuture-9999.0-empty-to-avoid-conflict-with-guava.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/log4j-1.2.17.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
metrics-core-3.2.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/mssql-
jdbc-6.2.1.jre7.jar:/usr/hdp/3.1.4.0-315/hbase/lib/netty-
all-4.0.52.Final.jar:/usr/hdp/3.1.4.0-315/hbase/lib/nimbus-jose-
jwt-4.41.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/objenesis-2.1.jar:/usr/
hdp/3.1.4.0-315/hbase/lib/okhttp-2.7.5.jar:/usr/hdp/3.1.4.0-315/hbase/
lib/okio-1.6.0.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
org.eclipse.jdt.core-3.8.2.v20130121.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
osgi-resource-locator-1.0.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
paranamer-2.3.jar:/usr/hdp/3.1.4.0-315/hbase/lib/phoenix-server.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/protobuf-java-2.5.0.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/ranger-hbase-plugin-shim-1.2.0.3.1.4.0-315.jar:/
usr/hdp/3.1.4.0-315/hbase/lib/ranger-plugin-
classloader-1.2.0.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/hbase/lib/
re2j-1.1.jar:/usr/hdp/3.1.4.0-315/hbase/lib/slf4j-api-1.7.25.jar:/usr/
hdp/3.1.4.0-315/hbase/lib/snappy-java-1.0.5.jar:/usr/hdp/3.1.4.0-315/
hbase/lib/spymemcached-2.12.2.jar:/usr/hdp/3.1.4.0-315/hbase/lib/stax2-
api-3.1.4.jar:/usr/hdp/3.1.4.0-315/hbase/lib/validation-
api-1.1.0.Final.jar:/usr/hdp/3.1.4.0-315/hbase/lib/woodstox-
core-5.0.3.jar:/usr/hdp/3.1.4.0-315/hbase/lib/xz-1.0.jar:/usr/hdp/
3.1.4.0-315/hbase/lib/zookeeper.jar:/usr/hdp/3.1.4.0-315/hadoop/conf:/
usr/hdp/3.1.4.0-315/hadoop/lib*/usr/hdp/3.1.4.0-315/hadoop/./*/usr/
hdp/3.1.4.0-315/hadoop-hdfs/./:/usr/hdp/3.1.4.0-315/hadoop-hdfs/lib*/:/
usr/hdp/3.1.4.0-315/hadoop-hdfs/./*/usr/hdp/3.1.4.0-315/hadoop-
mapreduce/lib*/usr/hdp/3.1.4.0-315/hadoop-mapreduce/./*/usr/hdp/
3.1.4.0-315/hadoop-yarn/./:/usr/hdp/3.1.4.0-315/hadoop-yarn/lib*/usr/
hdp/3.1.4.0-315/hadoop-yarn/./*/usr/hdp/3.1.4.0-315/tez*/usr/hdp/
3.1.4.0-315/tez/lib*/usr/hdp/3.1.4.0-315/tez/conf:/usr/hdp/
3.1.4.0-315/tez/conf_llap:/usr/hdp/3.1.4.0-315/tez/doc:/usr/hdp/
3.1.4.0-315/tez/hadoop-shim-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/
tez/hadoop-shim-2.8-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib:/
usr/hdp/3.1.4.0-315/tez/man:/usr/hdp/3.1.4.0-315/tez/tez-
api-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
common-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
dag-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
examples-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-history-
parser-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-javadoc-
tools-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-job-
analyzer-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
mapreduce-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-protobuf-
history-plugin-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
runtime-internals-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
runtime-library-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-

```

tests-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-yarn-timeline-
cache-plugin-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-yarn-
timeline-history-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/tez-
yarn-timeline-history-with-acls-0.9.1.3.1.4.0-315.jar:/usr/hdp/
3.1.4.0-315/tez/tez-yarn-timeline-history-with-
fs-0.9.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/ui:/usr/hdp/
3.1.4.0-315/tez/lib/async-http-client-1.9.40.jar:/usr/hdp/3.1.4.0-315/
tez/lib/commons-cli-1.2.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-
codec-1.4.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-
collections-3.2.2.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-
collections4-4.1.jar:/usr/hdp/3.1.4.0-315/tez/lib/commons-io-2.4.jar:/
usr/hdp/3.1.4.0-315/tez/lib/commons-lang-2.6.jar:/usr/hdp/3.1.4.0-315/
tez/lib/commons-math3-3.1.1.jar:/usr/hdp/3.1.4.0-315/tez/lib/gcs-
connector-1.9.10.3.1.4.0-315-shaded.jar:/usr/hdp/3.1.4.0-315/tez/lib/
guava-28.0-jre.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-
aws-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-
azure-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-azure-
datalake-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-hdfs-
client-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/hadoop-
mapreduce-client-common-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/
lib/hadoop-mapreduce-client-core-3.1.1.3.1.4.0-315.jar:/usr/hdp/
3.1.4.0-315/tez/lib/hadoop-yarn-server-timeline-
pluginstorage-3.1.1.3.1.4.0-315.jar:/usr/hdp/3.1.4.0-315/tez/lib/jersey-
client-1.19.jar:/usr/hdp/3.1.4.0-315/tez/lib/jersey-json-1.19.jar:/usr/
hdp/3.1.4.0-315/tez/lib/jettison-1.3.4.jar:/usr/hdp/3.1.4.0-315/tez/lib/
jetty-server-9.3.24.v20180605.jar:/usr/hdp/3.1.4.0-315/tez/lib/jetty-
util-9.3.24.v20180605.jar:/usr/hdp/3.1.4.0-315/tez/lib/
jsr305-3.0.0.jar:/usr/hdp/3.1.4.0-315/tez/lib/metrics-core-3.1.0.jar:/
usr/hdp/3.1.4.0-315/tez/lib/protobuf-java-2.5.0.jar:/usr/hdp/
3.1.4.0-315/tez/lib/RoaringBitmap-0.4.9.jar:/usr/hdp/3.1.4.0-315/tez/
lib/servlet-api-2.5.jar:/usr/hdp/3.1.4.0-315/tez/lib/slf4j-
api-1.7.10.jar:/usr/hdp/3.1.4.0-315/tez/lib/tez.tar.gz:/usr/hdp/
3.1.4.0-315/hadoop/conf:/usr/hdp/3.1.4.0-315/hadoop*/:/usr/hdp/
3.1.4.0-315/hadoop/lib*/:/usr/hdp/3.1.4.0-315/zookeeper/*::/conf:/lib/*'
-Djava.library.path=:/usr/hdp/3.1.4.0-315/hadoop/lib/native/Linux-
amd64-64:/usr/hdp/current/hadoop-client/lib/native/Linux-amd64-64:/usr/
hdp/3.1.4.0-315/hadoop/lib/native:/usr/hdp/3.1.4.0-315/hadoop/lib/
native/Linux-amd64-64:/usr/hdp/current/hadoop-client/lib/native/Linux-
amd64-64:/usr/hdp/3.1.4.0-315/hadoop/lib/native
org.apache.flume.node.Application --conf-file /home/admin/flm/flm.conf
--name StudentFlume
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/opt/apache-flume/lib/slf4j-
log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hadoop/lib/slf4j-
log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/phoenix/
phoenix-5.0.0.3.1.4.0-315-server.jar!/org/slf4j/impl/
StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
21/02/18 22:36:49 ERROR node.Application: A fatal error occurred while
running. Exception follows.
org.apache.commons.cli.ParseException: The specified configuration file
does not exist: /home/admin/flm/flm.conf

```

```
at org.apache.flume.node.Application.main(Application.java:342)
```

2. Создаем таблицы в hive

```
create external table lesson10_6.flm_logs (first string, second string,  
third string, fourth string, fifth string)
```

```
ROW FORMAT DELIMITED
```

```
FIELDS TERMINATED BY ','
```

```
STORED AS INPUTFORMAT
```

```
'org.apache.hadoop.mapred.SequenceFileInputFormat'
```

```
OUTPUTFORMAT
```

```
'org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat'
```

```
location '/user/centos/flm'
```

```
0: jdbc:hive2://bigdataanalytics-worker-0.nov> use lesson6_9;  
INFO : Compiling command(queryId=hive_20210218223912_b3952da7-  
d8a5-4e15-b0b2-4474e6f7f8d7): use lesson6_9  
INFO : Semantic Analysis Completed (retrial = false)  
INFO : Returning Hive schema: Schema(fieldSchemas:null,  
properties:null)  
INFO : Completed compiling  
command(queryId=hive_20210218223912_b3952da7-d8a5-4e15-  
b0b2-4474e6f7f8d7); Time taken: 0.01 seconds  
INFO : Executing command(queryId=hive_20210218223912_b3952da7-  
d8a5-4e15-b0b2-4474e6f7f8d7): use lesson6_9  
INFO : Starting task [Stage-0:DDL] in serial mode  
INFO : Completed executing  
command(queryId=hive_20210218223912_b3952da7-d8a5-4e15-  
b0b2-4474e6f7f8d7); Time taken: 0.011 seconds  
INFO : OK  
No rows affected (0.128 seconds)  
0: jdbc:hive2://bigdataanalytics-worker-0.nov> create external table  
lesson10_6.flm_logs (first string, second string, third string, fourth  
string, fifth string) ROW FORMAT DELIMITED FIELDS TERMINATED BY  
, 'STORED AS INPUTFORMAT  
'org.apache.hadoop.mapred.SequenceFileInputFormat' OUTPUTFORMAT  
'org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat' location '/  
user/centos/flm';  
INFO : Compiling  
command(queryId=hive_20210218223915_9eaff3a6-57ab-43c1-  
a3e4-4f27fb250bb5): create external table lesson10_6.flm_logs (first  
string, second string, third string, fourth string, fifth string) ROW  
FORMAT DELIMITED FIELDS TERMINATED BY ', 'STORED AS INPUTFORMAT  
'org.apache.hadoop.mapred.SequenceFileInputFormat' OUTPUTFORMAT  
'org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat' location '/  
user/centos/flm'  
INFO : Semantic Analysis Completed (retrial = false)  
INFO : Returning Hive schema: Schema(fieldSchemas:null,  
properties:null)
```

```

INFO : Completed compiling
command(queryId=hive_20210218223915_9eaff3a6-57ab-43c1-
a3e4-4f27fb250bb5); Time taken: 0.021 seconds
INFO : Executing
command(queryId=hive_20210218223915_9eaff3a6-57ab-43c1-
a3e4-4f27fb250bb5): create external table lesson10_6.flm_logs (first
string, second string, third string, fourth string, fifth string) ROW
FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS INPUTFORMAT
'org.apache.hadoop.mapred.SequenceFileInputFormat' OUTPUTFORMAT
'org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat' location '/
user/centos/flm'
INFO : Starting task [Stage-0:DDL] in serial mode
ERROR : FAILED: Execution Error, return code 1 from
org.apache.hadoop.hive.ql.exec.DDLTask.
MetaException(message:java.security.AccessControlException: Permission
denied: user=student1_13, access=WRITE, inode="/user":hdfs:hdfs:drwxr-
xr-x
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.check(FSPermi
ssionChecker.java:399)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:261)
    at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:193)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkPermission(FSDir
ectory.java:1857)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkPermission(FSDir
ectory.java:1841)
    at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkPathAccess(FSDir
ectory.java:1791)
    at
org.apache.hadoop.hdfs.server.namenode.FSNamesystem.checkAccess(FSNamesy
stem.java:7804)
    at
org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.checkAccess(Nam
eNodeRpcServer.java:2217)
    at
org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolServerSideTransl
atorPB.checkAccess(ClientNamenodeProtocolServerSideTranslatorPB.java:165
9)
    at
org.apache.hadoop.hdfs.protocol.proto.ClientNamenodeProtocolProtos$Clie
ntNamenodeProtocol$2.callBlockingMethod(ClientNamenodeProtocolProtos.java
)
    at
org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(P
rotobufRpcEngine.java:524)
    at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:1025)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:876)
    at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:822)
    at java.security.AccessController.doPrivileged(Native Method)

```



```

        at javax.security.auth.Subject.doAs(Subject.java:422)
        at
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformatio
n.java:1730)
        at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2682)
    )
INFO : Completed executing
command(queryId=hive_20210218223915_9eaff3a6-57ab-43c1-
a3e4-4f27fb250bb5); Time taken: 0.016 seconds
Error: Error while processing statement: FAILED: Execution Error, return
code 1 from org.apache.hadoop.hive ql.exec.DDLTask.
MetaException(message:java.security.AccessControlException: Permission
denied: user=student1_13, access=WRITE, inode="/user":hdfs:hdfs:drwxr-
xr-x
        at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.check(FSPermi
ssionChecker.java:399)
        at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:261)
        at
org.apache.hadoop.hdfs.server.namenode.FSPermissionChecker.checkPermissi
on(FSPermissionChecker.java:193)
        at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkPermission(FSDir
ectory.java:1857)
        at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkPermission(FSDir
ectory.java:1841)
        at
org.apache.hadoop.hdfs.server.namenode.FSDirectory.checkPathAccess(FSDir
ectory.java:1791)
        at
org.apache.hadoop.hdfs.server.namenode.FSNamesystem.checkAccess(FSNamesy
stem.java:7804)
        at
org.apache.hadoop.hdfs.server.namenode.NameNodeRpcServer.checkAccess(Nam
eNodeRpcServer.java:2217)
        at
org.apache.hadoop.hdfs.protocolPB.ClientNamenodeProtocolServerSideTransl
atorPB.checkAccess(ClientNamenodeProtocolServerSideTranslatorPB.java:165
9)
        at
org.apache.hadoop.hdfs.protocol.proto.ClientNamenodeProtocolProtos$Clien
tNamenodeProtocol$2.callBlockingMethod(ClientNamenodeProtocolProtos.java
)
        at
org.apache.hadoop.ipc.ProtobufRpcEngine$Server$ProtoBufRpcInvoker.call(P
rotobufRpcEngine.java:524)
        at org.apache.hadoop.ipc.RPC$Server.call(RPC.java:1025)
        at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:876)
        at org.apache.hadoop.ipc.Server$RpcCall.run(Server.java:822)
        at java.security.AccessController.doPrivileged(Native Method)
        at javax.security.auth.Subject.doAs(Subject.java:422)

```

```
        at
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformatio
n.java:1730)
        at org.apache.hadoop.ipc.Server$Handler.run(Server.java:2682)
) (state=08S01,code=1)
```

Практическая часть sqoop

1. Получаем список таблиц, которые созданы в postgres (нужно правильно указать ip, сейчас указан localhost)

```
sqoop list-tables --connect jdbc:postgresql://localhost/pg_db --username exporter
--password exporter_pass
```

2. Создаем таблицу в hive, куда собираемся выгружать данные CREATE TABLE `lesson6_9.character` (

```
`charid` string, `charname` string, `abbrev` string, `description` string,
`speechcount` int)
```

```
CREATE TABLE sell_prices
(store_id varchar, item_id varchar, wm_yr_wk varchar, sell_price varchar);
```

```
CREATE TABLE `sell_prices` ( `store_id` string, `item_id` string, `wm_yr_wk`
string, `sell_price` string)
```

3. Запускаем импорт:

```
sqoop import --connect jdbc:postgresql://localhost/pg_db --username exporter --
password exporter_pass --table sell_prices --hive-import --hive-database
student10_40 --hive-table sell_prices
```

```
CREATE TABLE `lesson6.character_p` (
```

© geekbrains.ru 4

```
`charid` string,
`charname` string,
`abbrev` string,
`description` string, `speechcount` int)
PARTITIONED BY (p_date string)
```

```
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe'
```

```
WITH SERDEPROPERTIES ( 'field.delim'='\u0001', 'line.delim'='\n',
'serialization.format'='\u0001')
```

```
STORED AS INPUTFORMAT 'org.apache.hadoop.mapred.TextInputFormat'
```

```
OUTPUTFORMAT
```

```
'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'
```

```
LOCATION 'hdfs://10.0.0.8:8020/user/hive/warehouse/character_p'
```

```
DATE=`date '+%Y%m%d'`
```

```
sqoop import --connect jdbc:postgresql://10.0.0.8/pg_db --username exporter --  
password exporter_pass --table character --target-dir /user/hive/warehouse/  
character_p/p_p_date=$DATE --delete-target-dir
```

```
hive -e 'msck REPAIR TABLE lesson6.character_p'
```

```
msck REPAIR TABLE lesson6.character_p
```

© geekbrains.ru 5

FLUME

1. Создаем shell-скрипт и сохраняем в файл heartbeat.sh

```
START_DATE=`date` COUNT=0
```

```
while [ true ]  
do
```

```
NOW_DATE=`date`
```

```
echo I live for $(( (`date -d "$NOW_DATE" +%s` - `date -d "$START_DATE" +  
%s`) )) seconds\; `(date -d "$START_DATE" +%Y-%m-%d:%H:%M:%S)\; `(date -d  
"$START_DATE" +%Y-%m-%d:%H:%M:%S)\; I did it $(( $COUNT + 1 )) times
```

```
COUNT=$(( $COUNT + 1 ))
```

```
sleep 10 Done
```

2. Создаем конфигурационный файл, правильно указываем пути до shell-скрипта и hdfs-папки # Naming the components on the current agent

```
StudentFlume.sources = Exec
```

```
StudentFlume.channels = MemChannel
```

```
StudentFlume.sinks = MySink
```

```
# Describing/Configuring the source StudentFlume.sources.Exec.type = exec
```

© geekbrains.ru 6

```
StudentFlume.sources.Exec.command = /home/centos/flm/heartbeat.sh
```

```
# Describing/Configuring the hdfs sink StudentFlume.sinks.MySink.type=hdfs
StudentFlume.sinks.MySink.hdfs.path=/user/centos/flm
StudentFlume.sinks.MySink.hdfs.fileSuffix=.log
```

```
# Describing/Configuring the channel StudentFlume.channels.MemChannel.type
= memory StudentFlume.channels.MemChannel.capacity=10000000
StudentFlume.channels.MemChannel.transactionCapacity = 100
```

```
# Bind the source and sink to the channel StudentFlume.sources.Exec.channels
= MemChannel StudentFlume.sinks.MySink.channel = MemChannel
```

3. Запускаем flume (нужно правильно указать пути до конфигурационной папки и конфигурационного файла)

```
/opt/apache-flume/bin/flume-ng agent --conf /home/user/flm/ --conf-file /home/
user/flm/flm.conf --name StudentFlume -Dflume.root.logger=INFO,console
```

4. Создаем таблицы в hive (найдите ошибку в скрипте, чтобы данные сразу раскладывались по колонкам, укажите правильно location)

```
create external table lesson10_6.flm_logs (first string, second string, third string,
fourth string, fifth string) ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
```

© geekbrains.ru 7

```
STORED AS INPUTFORMAT
'org.apache.hadoop.mapred.SequenceFileInputFormat' OUTPUTFORMAT
'org.apache.hadoop.hive ql.io.HiveSequenceFileOutputFormat' location '/user/
centos/flm';
```

5. Проверяем результат select * from flm_logs

© geekbrains.ru 8