

Домашнее задание

1. Подключить к Cassandra

```
[student1_13@bigdataanalytics-head-0 ~]$ /cassandra/bin/cqlsh 10.0.0.18
-bash: /cassandra/bin/cqlsh: No such file or directory
[student1_13@bigdataanalytics-head-0 ~]$ ssh 10.0.0.18
The authenticity of host '10.0.0.18 (10.0.0.18)' can't be established.
ECDSA key fingerprint is SHA256:FLX3eC5ZoVY/Ep1XEU1GV6VgbCtLwwMtBS2PDcYY0LM.
ECDSA key fingerprint is MD5:cd:54:0e:25:bb:b8:9d:04:92:57:89:af:db:12:cd:ed.
Are you sure you want to continue connecting (yes/no)? y
Please type 'yes' or 'no': yes
Warning: Permanently added '10.0.0.18' (ECDSA) to the list of known hosts.
[student1_13@bigdataanalytics-worker-2 ~]$ /cassandra/bin/cqlsh 10.0.0.18
Connection error: ('Unable to connect to any servers', {'10.0.0.18': error(111,
"Trying to connect to [('10.0.0.18', 9042)]. Last error: Connection refused"))}
```

2. Создать таблицы

3. Вставить записи

4. Изучить особенности работы where

• **Cassandra**

- Подключаемся к Cassandra на worker-2: `/cassandra/bin/cqlsh 10.0.0.18`
- Создаем пространство ключей: `CREATE KEYSPACE lesson7 WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 1 }`
- Создаем таблицу и вставляем значения: `CREATE TABLE animals (id int, name text, size text, primary key (id)); insert into animals (id, name, size) values (3, 'Deer', 'Big');`
- Проверяем как работает фильтрация: `select * from animals where id = 3 and name = '12321';`
- Сравниваем удаление и вставку пустого значения: `delete id from animals where id = 1;`
- `insert into animals (id, name, size) values (3, null, null);`

5. Подключиться к HBase

```
[student1_13@bigdataanalytics-worker-2 ~]$ hbase shell
SLF4J: Class path contains multiple SLF4J bindings.
```

```
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/phoenix/
phoenix-5.0.0.3.1.4.0-315-server.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/hdp/3.1.4.0-315/hadoop/lib/slf4j-
log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
For Reference, please visit: http://hbase.apache.org/2.0/book.html#shell
Version 2.0.2.3.1.4.0-315, r, Fri Aug 23 05:15:48 UTC 2019
Took 0.0016 seconds
hbase(main):001:0>
```

6. Создать таблицы и вставить значения

```
hbase(main):001:0> create_namespace 'lesson7'
```

```
ERROR: KeeperErrorCode = NoNode for /hbase-unsecure/master
```

Create namespace; pass namespace name,
and optionally a dictionary of namespace configuration.
Examples:

```
hbase> create_namespace 'ns1'
hbase> create_namespace 'ns1', {'PROPERTY_NAME'=>'PROPERTY_VALUE'}
```

Took 8.3662 seconds

```
hbase(main):002:0> create 'lesson7:animals', 'name', 'size'
```

```
ERROR: KeeperErrorCode = NoNode for /hbase-unsecure/master
```

Creates a table. Pass a table name, and a set of column family
specifications (at least one), and, optionally, table configuration.
Column specification can be a simple string (name), or a dictionary
(dictionaries are described below in main help output), necessarily
including NAME attribute.
Examples:

Create a table with namespace=ns1 and table qualifier=t1

```
hbase> create 'ns1:t1', {NAME => 'f1', VERSIONS => 5}
```

Create a table with namespace=default and table qualifier=t1

```
hbase> create 't1', {NAME => 'f1'}, {NAME => 'f2'}, {NAME => 'f3'}
hbase> # The above in shorthand would be the following:
hbase> create 't1', 'f1', 'f2', 'f3'
hbase> create 't1', {NAME => 'f1', VERSIONS => 1, TTL => 2592000, BLOCKCACHE
=> true}
hbase> create 't1', {NAME => 'f1', CONFIGURATION =>
{'hbase.hstore.blockingStoreFiles' => '10'}}
hbase> create 't1', {NAME => 'f1', IS_MOB => true, MOB_THRESHOLD => 1000000,
MOB_COMPACT_PARTITION_POLICY => 'weekly'}
```

Table configuration options can be put at the end.
Examples:

```
hbase> create 'ns1:t1', 'f1', SPLITS => ['10', '20', '30', '40']
hbase> create 't1', 'f1', SPLITS => ['10', '20', '30', '40']
```

```

hbase> create 't1', 'f1', SPLITS_FILE => 'splits.txt', OWNER => 'johndoe'
hbase> create 't1', {NAME => 'f1', VERSIONS => 5}, METADATA => { 'mykey' =>
'myvalue' }
hbase> # Optionally pre-split the table into NUMREGIONS, using
hbase> # SPLITALGO ("HexStringSplit", "UniformSplit" or classname)
hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit'}
hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit',
REGION_REPLICATION => 2, CONFIGURATION =>
{'hbase.hregion.scan.loadColumnFamiliesOnDemand' => 'true'}}
hbase> create 't1', {NAME => 'f1', DFS_REPLICATION => 1}

```

You can also keep around a reference to the created table:

```
hbase> t1 = create 't1', 'f1'
```

Which gives you a reference to the table named 't1', on which you can then call methods.

Took 8.1717 seconds

```
hbase(main):004:0> put 'lesson7:animals', '3', 'name', 'Deer'
```

```
^[[A
```

ERROR: Connection refused

Put a cell 'value' at specified table/row/column and optionally timestamp coordinates. To put a cell value into table 'ns1:t1' or 't1' at row 'r1' under column 'c1' marked with the time 'ts1', do:

```

hbase> put 'ns1:t1', 'r1', 'c1', 'value'
hbase> put 't1', 'r1', 'c1', 'value'
hbase> put 't1', 'r1', 'c1', 'value', ts1
hbase> put 't1', 'r1', 'c1', 'value', {ATTRIBUTES=>{'mykey'=>'myvalue'}}
hbase> put 't1', 'r1', 'c1', 'value', ts1, {ATTRIBUTES=>{'mykey'=>'myvalue'}}
hbase> put 't1', 'r1', 'c1', 'value', ts1, {VISIBILITY=>'PRIVATE|SECRET'}

```

The same commands also can be run on a table reference. Suppose you had a reference

t to table 't1', the corresponding command would be:

```
hbase> t.put 'r1', 'c1', 'value', ts1, {ATTRIBUTES=>{'mykey'=>'myvalue'}}
```

Took 139.7062 seconds

```
hbase(main):005:0>
```

```
hbase(main):006:0* put 'lesson7:animals', '3', 'name', 'Deer'
```

ERROR: Connection refused

Put a cell 'value' at specified table/row/column and optionally timestamp coordinates. To put a cell value into table 'ns1:t1' or 't1' at row 'r1' under column 'c1' marked with the time 'ts1', do:

```

hbase> put 'ns1:t1', 'r1', 'c1', 'value'
hbase> put 't1', 'r1', 'c1', 'value'
hbase> put 't1', 'r1', 'c1', 'value', ts1
hbase> put 't1', 'r1', 'c1', 'value', {ATTRIBUTES=>{'mykey'=>'myvalue'}}
hbase> put 't1', 'r1', 'c1', 'value', ts1, {ATTRIBUTES=>{'mykey'=>'myvalue'}}
hbase> put 't1', 'r1', 'c1', 'value', ts1, {VISIBILITY=>'PRIVATE|SECRET'}

```

The same commands also can be run on a table reference. Suppose you had a reference

t to table 't1', the corresponding command would be:

```
hbase> t.put 'r1', 'c1', 'value', ts1, {ATTRIBUTES=>{'mykey'=>'myvalue'}}
```

Took 138.7584 seconds

```
hbase(main):007:0> show namespaces
```

```
NameError: undefined local variable or method `namespaces' for main:Object
```

```
hbase(main):008:0> show namespace
```

```
NameError: undefined local variable or method `namespace' for main:Object
```

```
hbase(main):009:0> create 'lesson7:animals', {NAME => 'name', VERSIONS => 3}
```

```
ERROR: KeeperErrorCode = NoNode for /hbase-unsecure/master
```

Creates a table. Pass a table name, and a set of column family specifications (at least one), and, optionally, table configuration. Column specification can be a simple string (name), or a dictionary (dictionaries are described below in main help output), necessarily including NAME attribute.

Examples:

Create a table with namespace=ns1 and table qualifier=t1

```
hbase> create 'ns1:t1', {NAME => 'f1', VERSIONS => 5}
```

Create a table with namespace=default and table qualifier=t1

```
hbase> create 't1', {NAME => 'f1'}, {NAME => 'f2'}, {NAME => 'f3'}
```

```
hbase> # The above in shorthand would be the following:
```

```
hbase> create 't1', 'f1', 'f2', 'f3'
```

```
hbase> create 't1', {NAME => 'f1', VERSIONS => 1, TTL => 2592000, BLOCKCACHE => true}
```

```
hbase> create 't1', {NAME => 'f1', CONFIGURATION => {'hbase.hstore.blockingStoreFiles' => '10'}}
```

```
hbase> create 't1', {NAME => 'f1', IS_MOB => true, MOB_THRESHOLD => 1000000, MOB_COMPACT_PARTITION_POLICY => 'weekly'}
```

Table configuration options can be put at the end.

Examples:

```
hbase> create 'ns1:t1', 'f1', SPLITS => ['10', '20', '30', '40']
```

```
hbase> create 't1', 'f1', SPLITS => ['10', '20', '30', '40']
```

```
hbase> create 't1', 'f1', SPLITS_FILE => 'splits.txt', OWNER => 'johndoe'
```

```
hbase> create 't1', {NAME => 'f1', VERSIONS => 5}, METADATA => { 'mykey' => 'myvalue' }
```

```
hbase> # Optionally pre-split the table into NUMREGIONS, using
```

```
hbase> # SPLITALGO ("HexStringSplit", "UniformSplit" or classname)
```

```
hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit'}
```

```
hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit', REGION_REPLICATION => 2, CONFIGURATION =>
```

```
{'hbase.hregion.scan.loadColumnFamiliesOnDemand' => 'true'}}
```

```
hbase> create 't1', {NAME => 'f1', DFS_REPLICATION => 1}
```

You can also keep around a reference to the created table:

```
hbase> t1 = create 't1', 'f1'
```

Which gives you a reference to the table named 't1', on which you can then call methods.

Took 8.1909 seconds

```
hbase(main):001:0> create_namespace 'homework7'
```

ERROR: KeeperErrorCode = NoNode for /hbase-unsecure/master

Create namespace; pass namespace name,
and optionally a dictionary of namespace configuration.
Examples:

```
hbase> create_namespace 'ns1'  
hbase> create_namespace 'ns1', {'PROPERTY_NAME'=>'PROPERTY_VALUE'}
```

Took 15.3151 seconds

```
hbase(main):002:0> create_namespace 'hw7'
```

ERROR: KeeperErrorCode = NoNode for /hbase-unsecure/master

Create namespace; pass namespace name,
and optionally a dictionary of namespace configuration.
Examples:

```
hbase> create_namespace 'ns1'  
hbase> create_namespace 'ns1', {'PROPERTY_NAME'=>'PROPERTY_VALUE'}
```

Took 18.0665 seconds

7. Изучить особенности хранения данных

• **HBase**

• Подключаемся к HBase *hbase shell*

```
create_namespace 'lesson7'  
create 'lesson7:animals', 'name', 'size'
```

• Вставляем значения:

```
put 'lesson7:animals', '3', 'name', 'Deer' put 'lesson7:animals', '3', 'size', 'Big'  
put 'lesson7:animals', '5', 'name', 'Snake' put 'lesson7:animals', '3', 'name', 'Doe'
```

• Удаляем значение:

```
delete 'lesson7:animals', '5'
```

• Делаем запрос к созданной таблице: *get 'lesson7:animals', '5'*